

# 1 Problem

Existing Frameworks are statically allocating resources to jobs

- can't maintain SLOs while maintaining full utilization and efficiency

- 1 solution killing tasks - high overheads

- 1 solution only use 70% of resources

—— Utilization: using all resources; Efficiency: not doing tasks

# 2 not working solutions

make task size smaller - more disk seeks and more overhead at scheduler

uniform size tasks - doesn't work because  $\frac{\text{input amount}}{\text{task time}}$  is not consistent

OS checkpoints, suspend and save state in memory - too much space needed, HDD no option too slow etc.

# 3 Solution

Amoeba - Prototype

Smallest granularity is task - make it more granular - split up tasks

Mapper gets set of records as input (a list of words); every possible sublist is a admissible partitioning of the data (the blocks to work on can consist of any number of words); you can kill a task after every record and pretend you made the blocks that way A Checkpoint is the next key where to continue