

1 Parallel computing

Doing things concurrently \cong more complex than sequential programming. Central questions:

- BASICS: How do we split up tasks?
- BASICS: How do the different parts communicate?
- PERFORMANCE: How much time does it take to run the application? (User's question)
- PERFORMANCE: Is the system fully utilized? (Adim's question)

A few important terms are:

Computational Model: Model that can be solved with a computer

Model of Computation: The description of how to solve a Computational Model with a computer Typ-

Concurrent: Parallel / At the same time

ical steps in development are:

Specification \rightarrow Design \rightarrow Implementation \rightarrow Adaption to system \rightarrow Adaption to Machine

1.1 Invisible Parallelism

This is parallelism that we don't see while programming like:

- In compiler for the hardware and the OS
- Pipelining
- On the Bit-/Instruction-/Memory-Level

1.2 Amdahl's Law

Describes how much a program can be sped up:

$$S = \frac{1}{s + (1-s)/p} \geq \frac{1}{s}$$

S: speed up

s: time used for sequential part

1-s: time used for parallel part

p: number of processors

1.3 Granularity

- Splitting up processes into pieces
- How can they communicate? How can I combine them to run on 1 processor?
- How small can and should I make them?

1.4 Data Locality

- How to minimize data retrieval time?
- How to keep data as local as possible?

1.5 Load Imbalance

- Trying to balance load across processors
- Balancing the load of regular and irregular processes

2 Parallel Algorithms

How do you identify parallelism?

- Look up algorithms
- Look at code
- Look at data structures

2.1 Design

1. Decompose (Machine Dependent)
 - What level of parallelism? (static / dynamic)
 - How many processors?
 - How big is the problem?
2. Clustering (Machine Independent)
 - How to bring tasks together?
3. Implement (Machine Dependent)
4. Map work to processes (Machine Dependent)
 - Doing it static / dynamic?
 - Being aware of data locality and possible load imbalance.

2.2 Memory

There are different kinds of memory:

- Distributed: Programmer has to manually communicate data. This is most common.
- Shared: Uniform location of data → easier for programmer. The bus can get overloaded here.
- Virtual Shared: Is distributed but pretends to be shared.

2.3 Solving Concepts

1. Farmer Worker
 - Farmer organizes and schedules
 - Workers do work
 - Farmer can push work to workers (Push-Model)
 - Workers can pull work from farmer (Pull-Model)
2. Divide & Conquer
 - Recursive split up (for example)
 - Hardest part is the Administration
 - First divide problem then conquer subparts
3. Data Parallelism
 - When data is splitable
 - Different nodes work on different parts of data
4. Task Parallelism
 - Different tasks for different nodes
 - Task can start when input is available (data dependent)
 - Task can start when its output is needed (demand-flow)
5. Bulk Synchronisation
 - Every part waits at a certain point for every other part and continues together

2.4 System Level Communication

With Shared Memory

- Mutual exclusion per data (Only one can access)
- (Un)locking mechanisms put in place
- Barriers put in place where processes wait

With Distributed Memory

- Communication has to happen
- After sending data, it is duplicate → which is the original?
- General structure of send command: SEND(Buffer*, element-count, destination)
- General structure of receive command: RECEIVE(Buffer*, element-count, owner)